

# **DBMS - 2**

downloaded from  
[csitauthority.blogspot.com](http://csitauthority.blogspot.com)



31-10-14

## Referential Integrity in SQL

1. Create table customer (

customer\_name char(100),  
customer\_street char(100),  
customer\_city char(100),  
primary key (customer\_name)

);

2. create table branch (

branch\_name char(100),  
branch\_city char(100),  
assets int,  
primary key (branch\_name)

);

3. create table account (

account\_number char(100),

branch\_name char(100),

balance int,

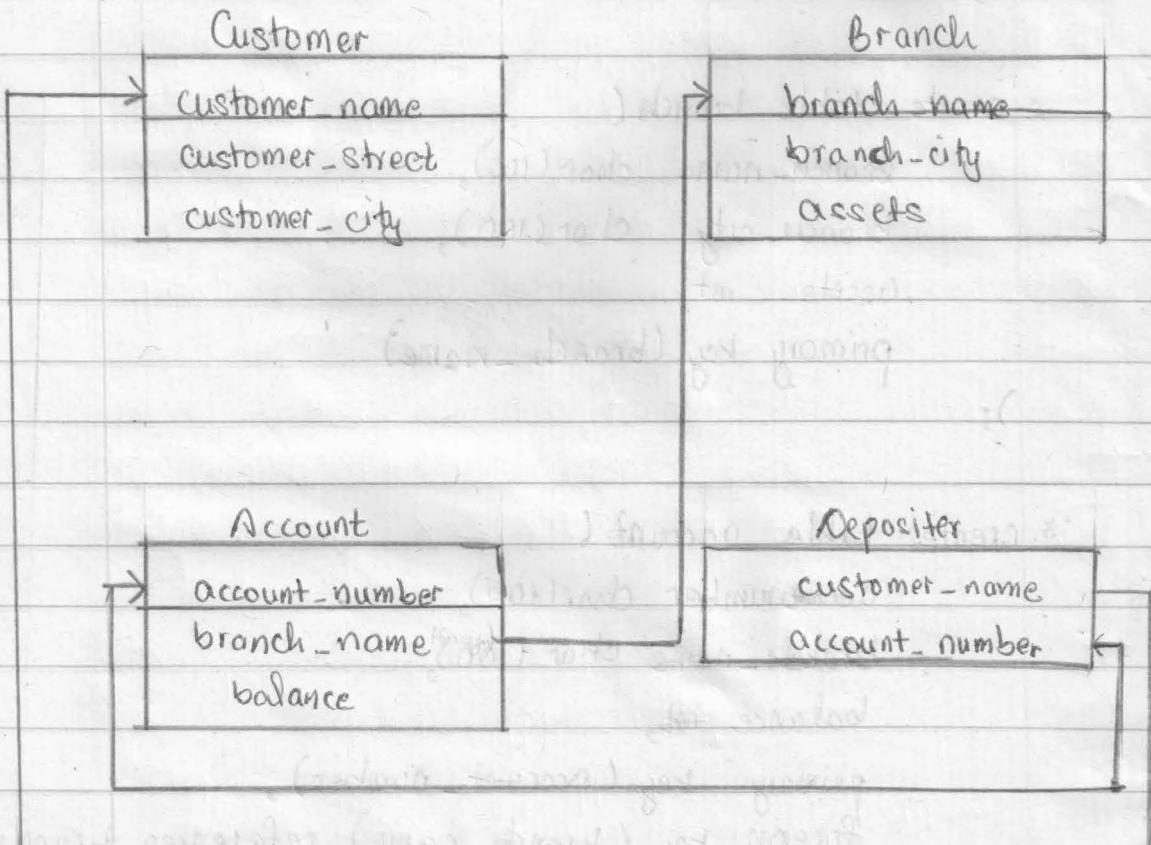
primary key (account\_number),

foreign key (branch\_name) references branch

);

1. Create table depositor (

customer\_name char(100),  
 account\_number char(100),  
 primary key (customer\_name, account\_number),  
 foreign key (account\_number) references account,  
 foreign key (customer\_name) references customer  
 );





## Assertions

An assertion is a predicate expressing a condition ~~with~~ <sup>wish</sup> which the database to always satisfy. Domain constraints, functional dependency and referential integrity are special forms of assertion.

If a constraint cannot be expressed in these forms, we use an assertion.

Example: sum of loan amounts for each branch is less than the sum of all accounts' balances at the branch.

Every loan customer keeps a minimum of thousand dollars in an account.

General syntax for creating assertion in SQL is  
create assertion <assertion-name> check <predicates>

1. create assertion sumconstraint check

(not exist (select \* from branch

where (select sum(amount) from loan

where loan.branch\_name = branch.branch\_name)

>= (select sum(amount) from account

where loan.branch\_name = branch.branch\_name))

## Fundamental Dependencies

Fundamental dependencies are constraints on the set of legal relation. It defines attributes of relation, how they are related to each-other.

It determines unique value for certain set of attributes to the value for another set of attributes i.e. functional dependency is a generalisation of the notation of key.

Fundamental dependencies are interrelationships among attributes of a relation.

"For a given relation R with attributes X and Y, Y is said to be fully functional dependent on X, if given value for each  $x \in X$  uniquely determines the value of the attribute in Y. X is called determinant of the functional dependency (F.D) and functional dependency is denoted by  $X \rightarrow Y$ ".

Example: Consider a relation supplier

supplier(supplier\_id, sname, status, city)

- supplier.supplier\_id  $\rightarrow$  supplier.sname
  - supplier.supplier\_id  $\rightarrow$  supplier.status
  - supplier.supplier\_id  $\rightarrow$  supplier.city
- $\left. \begin{matrix} \end{matrix} \right\} \text{supplier\_id} \rightarrow \text{sname}$

There are four types of functional dependencies:

- i. Full functional dependency
- ii. Partial functional dependency
- iii. Transitive dependency
- iv. Trivial functional dependency

10-11-14 Assignment #4

### Relational Calculus

- Tuple Relational Calculus
  - Well formed formula
  - Safety of Expressions
  - Equivalent expressions
  - Examples of queries in Tuple Relations (Basic Operations)
- Domain Relational calculus (Basic Operations Only)

## Trigger

Trigger is a statement that is automatically executed by the system as a side-effect of a modification to the database. While writing a trigger, we must specify:

- conditions under which the trigger is executed.
- action to be taken when trigger executes.

Triggers are useful mechanism to perform certain task automatically when certain condition is meet. Sometimes trigger is also called rule or action rule.

### Basic Syntax for trigger:

Create or Replace Trigger <Trigger\_name>  
 {Before, After}

{Insert | Delete | Update [of column, ...]}

ON <Table\_name>

[Referencing Old as <old>, New as <new>]

Declare

variable declaration ;

Begin

End;

Example:

Suppose that instead of allowing negative account balances, the bank deals with overdrafts by:

- setting the account balance to zero.
- creating a loan in the amount of the overdraft providing same loan number as account number of the overdrawn account.

create trigger overdraft\_trigger

after update on account

referencing new row as nrow

for each row

when nrow.balance < 0

begin

insert into borrower

(select customer\_name, account\_number

from depositor

where nrow.account\_number,

nrow.branch\_name, nrow.balance);

update account

set balance = 0

where account.account\_number

= nrow.account\_number

end;

end;

Why triggers can be hard to understand?

In an active database system, when the DBMS is about to execute a statement that modifies the database, it checks if some triggers are activated by the statement. If so, the DBMS processes the trigger by evaluating its conditional condition part, and then (if the condition evaluates to true) is executing its action part.

If a statement activates more than one triggers, the DBMS typically processes all of them in one arbitrary order. An important point is that the execution of the action part of a trigger could end turn activate another trigger. The execution of the action part of a trigger could again activate another trigger. Such triggers are called recursive triggers. The potential for such chain activation, and the unpredictable order in which a DBMS process activated triggers can make it difficult to understand the effect of a collection of triggers.

## Constraints v/s Triggers (Do yourself)

## Other uses of triggers

- i. Trigger can alert users to unusual events (as reflected in updates to the database).  
For eg: we may want to check whether a customer placing an order has made enough purchase in the last month to qualify for an additional discount. If so, the sales clerk must be informed so that he can tell the customer and possibly generate additional sales.
- ii. We can rely this information by using a trigger that checks recent purchase and print a message if the customer qualifies for the discount.
- iii. Trigger can generate a log of events to support auditing and security checks.
- iv. We can use triggers to gather statistics on table accessed and modifications.

## String Operations

like - operator

i. Percent (%) : matches any substrings

ii. Underscore (\_) : matches any character

Eg: 'I %' matches any string beginning with I

'ABC%'

'% I %'

'%pur'

'I % I'

'%.I.%' - matches any string containing 'I' as substring

like '%\_\_%' matches any string of exactly three characters.

like '%\_\_\_%' matches any string of <sup>at least</sup> exactly three characters

List all customers whose name starts with S.

Select customer\_name from customer

where customer\_name like 'S%';

24-11-14

24-11-14

## Assignment # 5

1. Trivial and non-trivial full dependencies, closure of a set of functional dependencies.  
Attribute closure FDs, irreducible set of FDs
2. Transitivity, reflexivity and augmentation properties of FDs.

1-12-14 Job assignment # 7

### # Aggregate functions

min, max, sum, avg, count, count\*

1. How many tuples are stored in the relation employee.
2. How many different job titles are stored in the job title employee?
3. List the minimum and maximum salary.
4. Sum of all salaries of employees working in dept 2.

## Anomalies

An anomaly is simply an error or inconsistency in the database. A poorly designed database runs the risk of introducing numerous anomalies.

There are three types of anomalies:

### i. Insertion anomalies

An anomaly that occurs during the insertion of a record. It is the inability to add data to the database due to absence of other data.

### ii. Deletion anomalies

An anomaly that occurs during the deletion of a record. It is the unintended loss of data due to deletion of other data.

### iii. Update anomalies

An anomaly that occurs during the updating of a record. For example, updating a description's column for a single part in an inventory database require you to make a change to thousands of rows.

## Pitfall of relational model:

The main goal of relational database is:  
i. to create / find good collection of relational schema.  
such that database should allow us to store  
data / information without unnecessary redundancy  
and should also allow to retrieve information  
easily, i.e. the goal of relational database design  
should be concentrated

- to avoid redundant data from database
- to ensure that relationship among attributes  
are represented.
- to facilitate the checking of updates for the  
violation of data integrity constraints

ii. A bad relational database design may lead to

- repetition of information i.e. it leads data  
redundancy in database, so obviously it  
requires much space
- inability to represent certain information

Example: Consider a relational schema

branch-loan = (branch-name, branch-city, assets,  
customer-name, loan-no, amount)

branch-name	branch-city	assets	customer-name	loan-number	amount
Kathmandu	Baneshwor	25000	Rohan	L-15	3000
Kalitpur	Patan	19000	Mohan	L-17	5000
Kathmandu	Baneshwor	25000	Raju	L-19	10000
Pokhara	Prithvinagar	17000	Manoj	L-10	7000
Kalitpur	Patan	19000	Swikar	L-30	9000

### Redundancy

- i. Data for branch-name, branch-city, assets are repeated for each loan that is provided by bank
- ii. Storing information several times leads wastage of storage space/time.
- iii. Data redundancy leads problem in insertion, deletion and update

**i) Insertion Problem:** We cannot store information about a branch without loan information, we can use NULL values for loan information but they are difficult to handle.

### ii) Deletion problem

In this example, if we delete the information of Manoj (i.e. delete from branch-loan where

customer\_name='Manoj'), we cannot obtain the information of Pokhara branch i.e. we don't know branch\_city of Pokhara, total assets of Pokhara branch, etc.

### iii. Update Problem

Since data are duplicated, so multiple copies of same fact need to update while updating one. It increases the possibility of data inconsistency when we made update in one copy, there is possibility that only some of the multiple copies are updated but not all, which lead data / database in inconsistent state.

### Decomposition

The idea of decomposition is breakdown of large and complicated relation in number of simple and small relations which minimizes data redundancy. It can be considered principle to solve the relational model problem.

Lossless join decomposition - Do yourself

~~10-12 marks~~

# NORMALIZATION

## Concept of normalization

Relational database is a collection of tables. The database must be carefully designed in order to get full advantage. By designing database tables carefully, we can save space, minimize duplicate data, protect data consistency as well as provide faster transaction by sending less data.

Normalization is the answer to achieve these benefits. It is a process of splitting tables to reduce data redundancy (duplicate data) and establishing relation between tables. Normalization provide flexibility, data consistency (or data integrity) and avoid anomalies while inserting, deleting, updating data.

The following figure illustrates the unnormalized and normalized database.

Last name
First name
Birthday
Home phone
Home city
Home Street
Date of call
Call description



Table: Contact

Contact\_id

Last name

First name

Birthday

Home phone

Home city

Home street

Table: Note

Contact\_id

Date of call

Call description

Call type id

Fig. Normalized form

Table: Calltype

Call type id

Call description

Fig. Next stage of normalized table

Needs of Normalization - Do yourself.

## Objectives of Normalization

- to free a collection of relation from undesirable insertion, update and deletion.
- to move the relational model more informative to use.
- to increase the lifetime of application programs.

## Properties of Normalization Relation after normalization

- no data value should be duplicate in different row or tuple.
- a value must be satisfied for every attribute in a tuple.
- important information should not be accidentally lost.
- when new data/record are inserted to relation, other relation in a database should not be affected.

## Forms of normalization

The breaking down of a table may undergo series of stages called as normal form or just an NF in short. There are 6 normal forms so far, which are widely used. The higher level of normalization can not be achieved unless previous levels have been satisfied. For example, to be in second normal form, it should already be in first normal form.

- First normal form (1NF)
- Second normal form (2NF)
- Third normal form (3NF)
- Boyce Codd Normal form (BCNF)
- Fourth normal form (4NF)
- Fifth normal form (5NF)

### First Normal Form (1NF)

A table is said to be first normal form if it has no repeating groups i.e. for each cell in a table, there can be only one value. If a group of five items repeat, it should be split into a new table.

Let us consider an example of billingsystem of a store.

Sales_order no
Date
Customer No
Customer Name
Customer Address
Item No
Description
Quantity
Unit price

Fig. A  
Unnormalized  
Table

Table: Customer

Sales_order_no
Date
Customer_no
Customer_name
Customer_address

Table : Sales

Sales_order_no
ItemNo
Description
Quantity
Unit price

Fig. Table in 1NF

The fields date, customer\_no, customer\_name, customer\_address are repeating fields and the fields itemno, description, quantity, unit price are non-repeating fields. If all fields are kept in a single table, header data are repeated for every line in sales order i.e. to keep record for each item's information on customer should also be entered each time.

In 1NF, repeating groups are separated into new table as shown in figure above.

These 2 tables can be linked together by a common field - sales\_order\_no. Now, these 2 tables are in 1NF.

## Second Normal Form (2NF)

A table is said to be in second normal form if it is already in first normal form and if every non-key columns depend on the entire key.

Let's discuss about sales table.

Table: Sales

Sales order no
Item No
Description
Quantity
Unit price

Here, the key field is sales order no. The field description depends on Item No rather than sales order no. However, quantity and unit price are not dependent on item no as they may be different for different sales order. In this case, the table is split again. The column that depend on the key are kept in one table and rest on another table as follows.

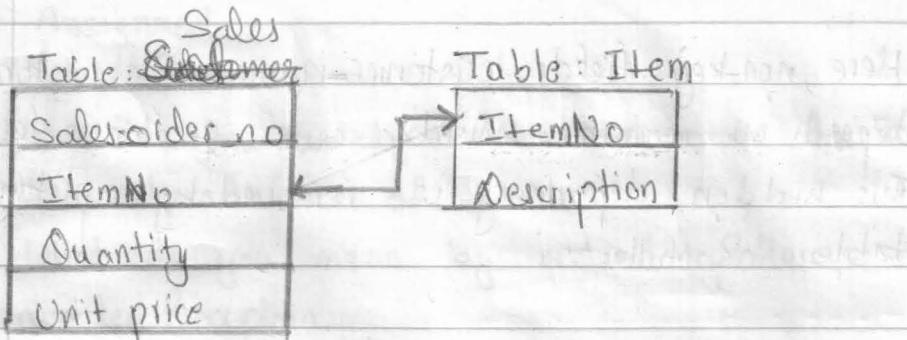


Fig. Table in 2NF

Now, these two tables are in 2NF and they can be linked together by common field ItemNo.

### Third Normal Form (3NF)

A table is said to be in 3NF if it is already in 2NF and if non-key columns are not dependent on each other. There should not be any hidden dependencies among non-key columns. Let's discuss about customer table.

Table : Customer

Sales-order_no
ItemNo
Quantity
Unit price

Table : Customer

Sales-order_no
Date
Customer_no.
Customer_name
Customer_address

Here, non-key fields customer\_name and customer\_address depends on customer\_number but not on sales\_order\_no. This hidden dependency is removed by splitting the table as follows:

Table : Customer	Table : Sales order
<u>Customer no</u>	<u>Sales order no</u>
<u>Customer name</u>	<u>Date</u>
<u>Customer address</u>	<u>Customer no</u>

Fig. Table in 3NF

Now, the final normalized tables after third normal form will be as follows:

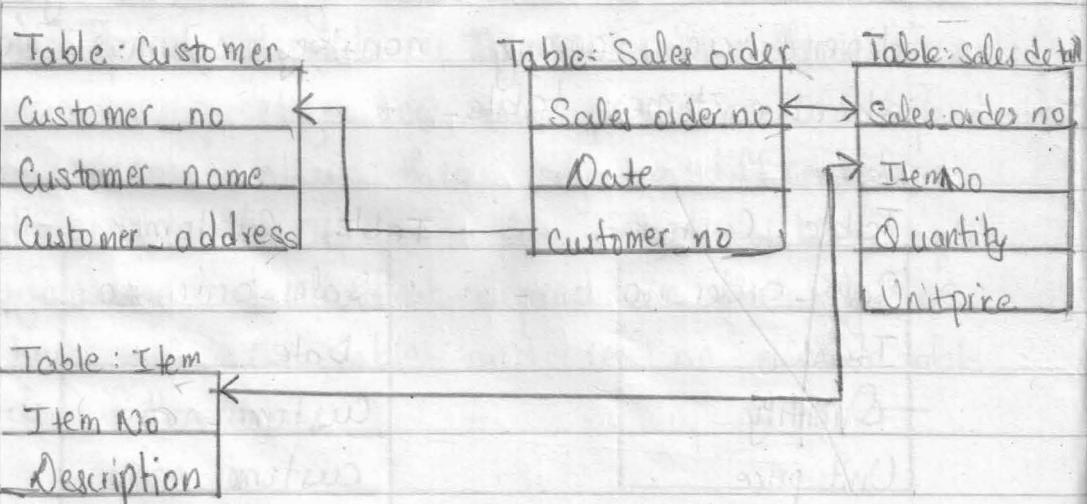


Fig. Final normalized table

8-12-14

Assignment  
Lab # 7

1. What do you mean by relational database design? List out the characteristics of RDBMS.
2. What do you mean by redundancy? How can this be avoided? Explain.
3. Define term anomalies. Explain BCNF in detail.
4. Distinguish the differences between 2<sup>nd</sup> and 3<sup>rd</sup> NF.
5. What do you mean by functional dependency. What role does it play in normalization? Explain.
6. What is transitive dependency? How does it differ from full functional dependency? Explain.
7. Why is concurrency control needed? Explain lost update, inconsistent retrieval and uncommitted dependency anomalies.

# TRANSACTION MANAGEMENT

## Recovery

Recovery in a database system means, primarily, recovering the database itself, i.e. restoring the database to a state that is known to be correct (or, rather, consistent) after some failure have rendered the current state inconsistent, or at least suspect.

## Commit

The commit operation signals successful end of transaction, it tells the transaction manager that a logical unit of work has been successfully completed, the database is in a consistent state again and all of the updates made by that unit of work can now be "committed" or made permanent.

### Rollback :

The rollback operation signals unsuccessful end of transaction, it tells the transaction manager that something has gone wrong, the database might be in an inconsistent state, and all of the updates made by the logical unit of work so far must be "rolled back" or undone.

### Transaction

Transaction begins with successful execution of a begin transaction statement and it ends with successful execution of either a commit or a rollback statement.

A commit point thus corresponds to the end of a logical unit of work and hence to a point at which the database is or should be in a consistent state. Rollback, by contrast, rolls the database back to the state it was in at begin transaction, which effectively means back to the previous commit point.

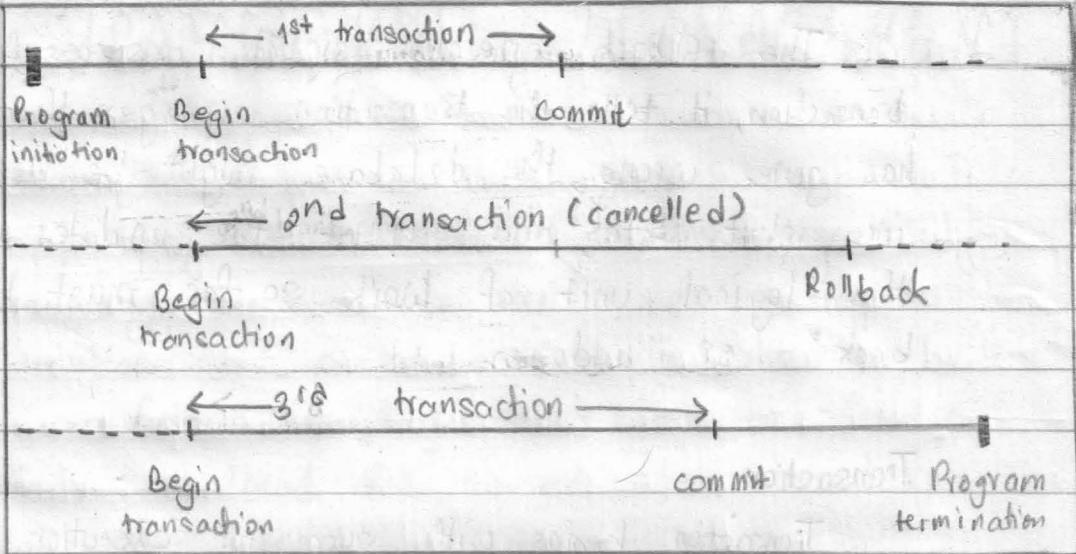


Fig. Program collection is a sequence of transactions

When a commit point is established

- all updates made by the executing program, since the previous commit points are committed, i.e. they are made permanent. Once committed, an update is guaranteed, never to be undone
- all database positioning is lost and all tuple locks are released. Database positioning here refers to the idea that at any given time, an executing program will typically have addressability to certain tuples

## The ACID properties

- Atomicity      All or nothing
  - Consistency
  - Isolation
  - Durability
- 
- Transactions are **atomic** (all or nothing)
  - Transactions preserve database **consistency**, i.e. a transaction transforms a consistent state of a database into another consistent state, without necessarily preserving consistency at all intermediate point

## System Recovery

The system mostly prepared to recover, not only from purely local failures such as the occurrence of an overflow condition within an individual transaction but also from global failures such as power outage.

A local failure, affects only the transaction in which the failure has actually occurred, such failure has

A global failure, affects all of the transaction in progress at the time of the failure, and hence the significant system wide implementation wide implication.

Failures fall into two broad categories:

- i. System failures (E.g. power outage), which affect all transactions currently in progress but do not physically damage the ~~system~~ database. A system failure is sometimes called soft crash.
- ii. Media failures (E.g. head crash on the disk), which do cause damage to the database, or to some portion of it, and affect atleast those transaction currently using that portion. A media failure is sometimes called a hard crash.

## Media Recovery

A media failure is a failure such as a disk head crash or a disk-controller failure in which some portion of the database has been physically destroyed. Recovery from such a failure basically involves reloading (or restoring). The database from a backup copied and then using the log both active and ~~archived~~ portion.

The need to be able to perform media recovery ~~implies~~ comprise the need for a dump/restore (or unload/reload) utility.

After a media failure, the restore portion of the utility is used to create the database from a

~~Utility~~: specified backup copy.

## Two-phase commit.

- Commitment

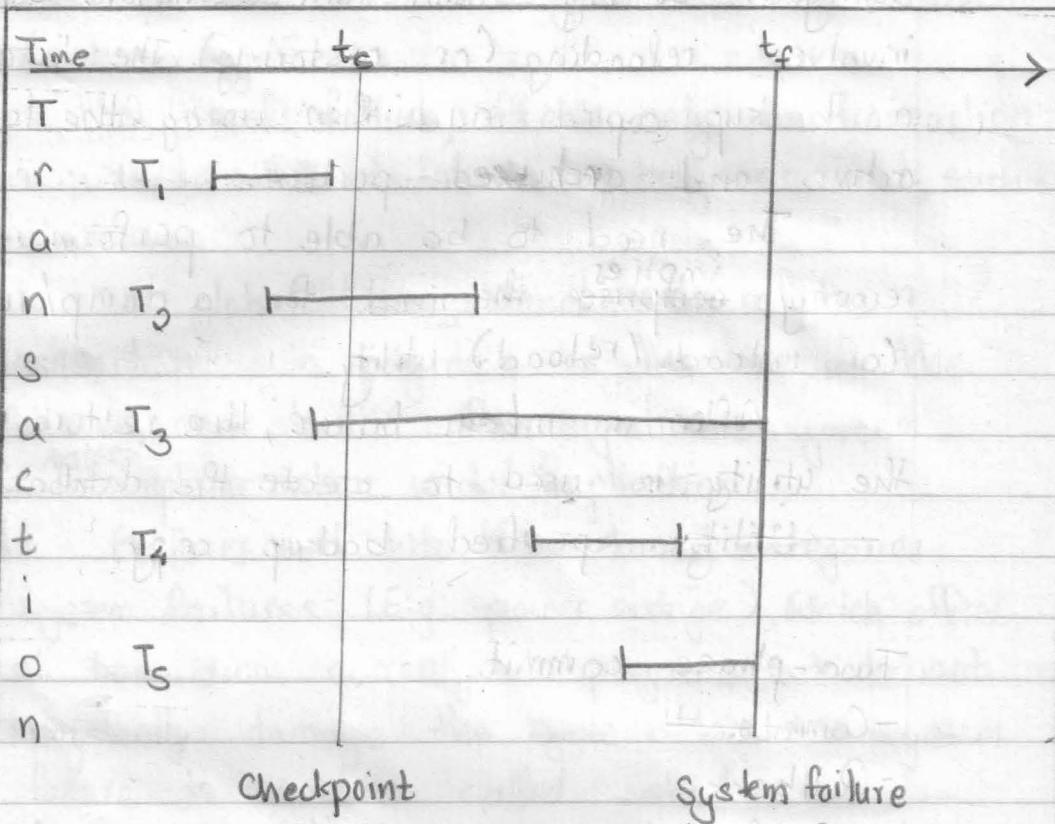
- Rollback

A very important elaboration on the basic commit/rollback concept called two phase commit.

Two-phase commit is important whenever a given transaction can interact with several independent resource manager, each managing its own set of

recoverable resources and maintaining its own recovery log

### Transaction categories



## 5 transaction categories

- a system failure has occurred at time  $t_f$ .
- the most recent checkpoint prior to time  $t_f$  was taken at time  $t_c$ .
- transaction type  $T_1$  completed (successfully) prior to time  $t_c$ .
- transaction type  $T_2$  started prior to time  $t_c$  and completed after time  $t_c$  and before time  $t_f$ .
- transaction type  $T_3$  also started prior to time  $t_c$  but did not complete by time  $t_f$ .
- transaction type  $T_4$  started after time  $t_c$  and completed before time  $t_f$ .
- finally, transaction type  $T_5$  also started after time  $t_c$  but did not complete by time  $t_f$ .

## Concurrency

The term concurrency refers to the ~~cracke~~ fact that DBMSs typically allow many transactions to access the same database at the same time, and in such a system as is well-known, some kind of concurrency control mechanism is needed to ensure that concurrent transactions do not interfere with each other.

## Three Concurrency Problems

- i the lost update problem
- ii the uncommitted dependency problem
- iii the inconsistent analysis problem

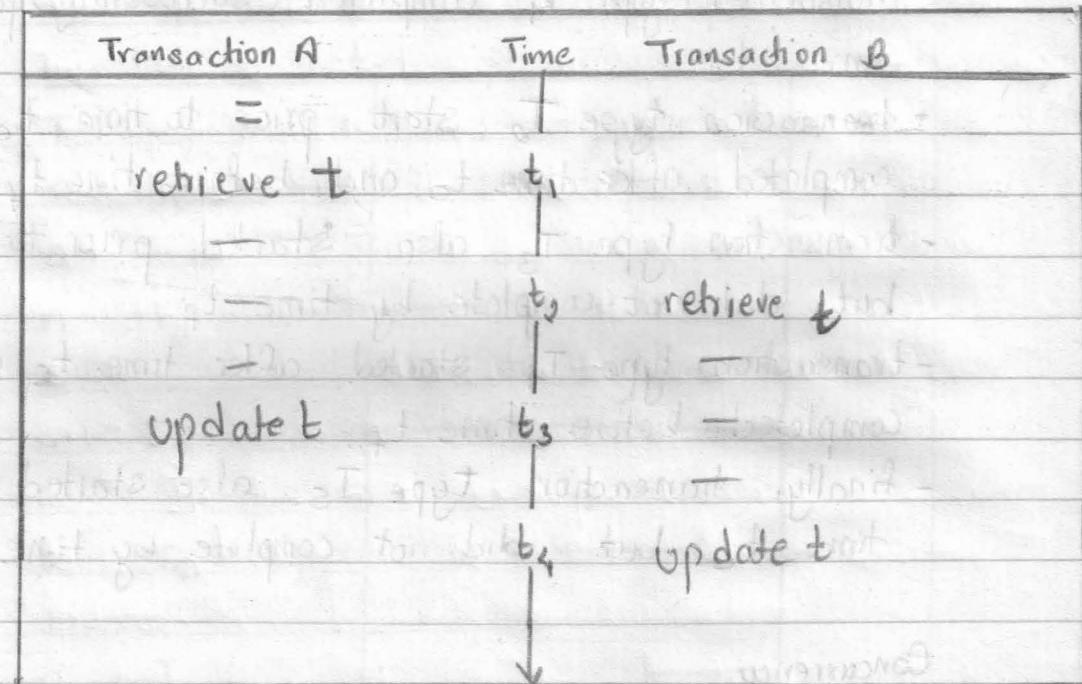


Fig. Transaction A lose an update at time t

Transaction A retrieves some tuple  $t$  at time  $t_1$ ; transaction B retrieves that same tuple  $t$  at time  $t_2$ .

Transaction A updates the tuple (on the basis of the values seen at time  $t_1$ ) at time  $t_3$  and

transaction B updates the same tuple at  $t_4$ .

Transaction A's update is lost at time  $t_4$ , because transaction B overwrites it without even looking at it.

## Lab #8

### Database recovery

- Purpose of database recovery
- Types of failure
- Storage hierarchy
- Buffer management
- transaction log

## The uncommitted Dependency Problem

The uncommitted dependency problem arises if one transaction is allowed to retrieve, worse, update - a tuple that has been updated by another transaction but not yet committed by that other transaction. For if it has not yet been committed, there is always a possibility that it never will be committed but will be rolled back instead - in which case the first transaction will have seen some data that now no longer exist (and in a sense never did exist).

In the first example (which is shown below), transaction A sees an uncommitted update (also called an uncommitted change) at time  $t_2$ . That update is then undone at time  $t_3$ .

Transaction A	Time	Transaction B
—		—
—	$t_1$	Update t
—		—
Retrieve t	$t_2$	—
—		—
—	$t_3$	Rollback
	↓	

Fig.a Transaction A becomes dependent on an uncommitted change at time  $t_2$ .

Transaction A	Time	Transaction B
—		—
—	$t_1$	Update t
—		—
Update t	$t_2$	—
—		—
—	$t_3$	Roll back
	↓	

Fig.b Transaction A updates an uncommitted change at time  $t_2$ , and loses that update at time  $t_3$ .

## The inconsistent analysis problem

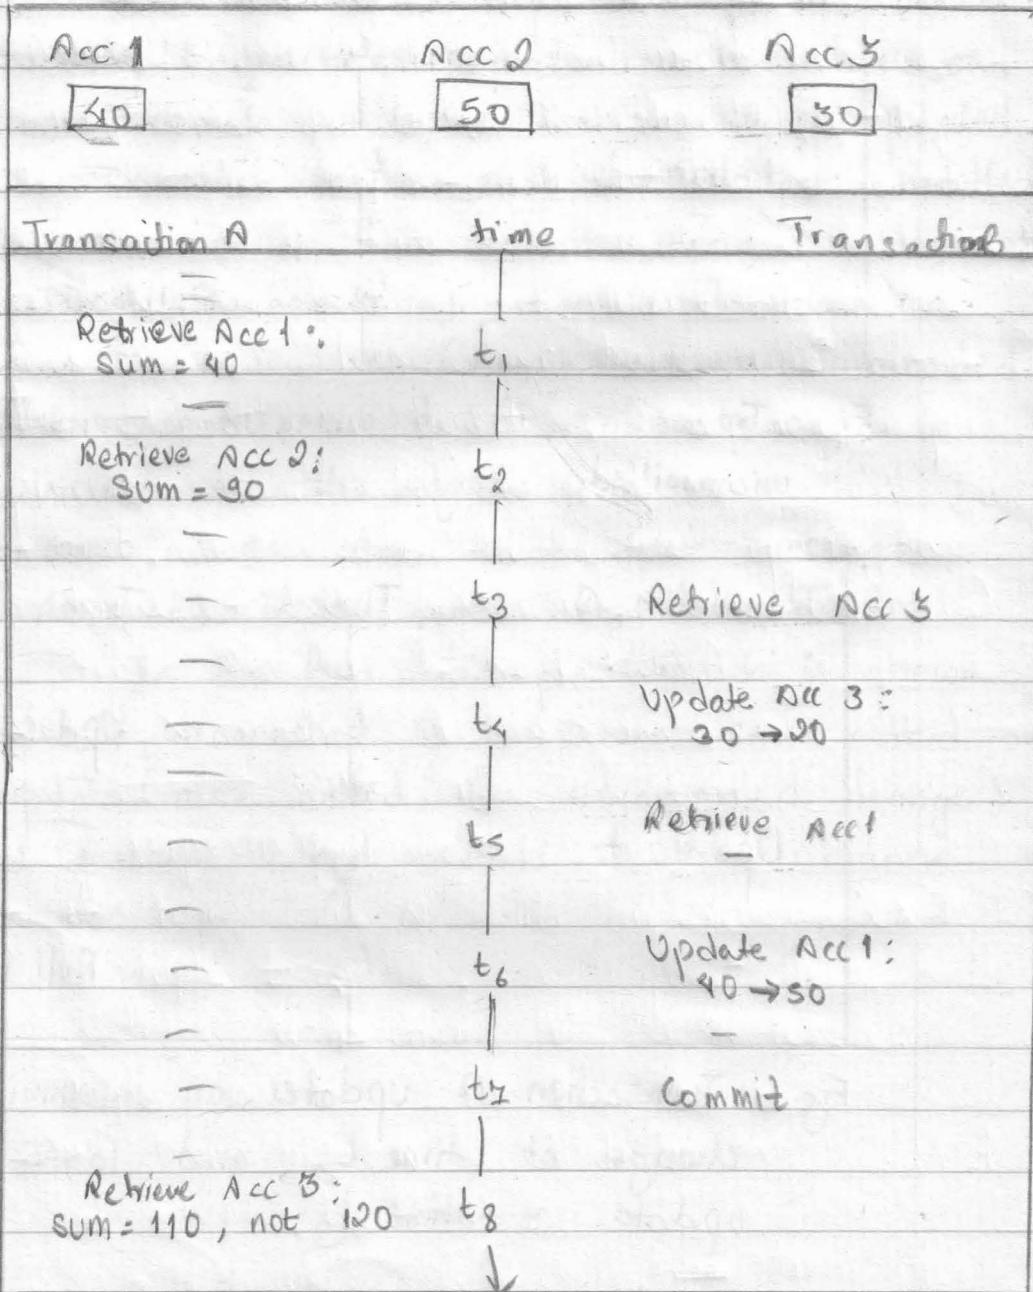


Fig. transaction A performs an inconsistent analysis

Consider the figure, which shows two transactions A and B operating on account tuples: transaction A is summing account balances, transaction B is transferring an amount 10 from account 3 to account 1. The result produced by A, 110, is obviously incorrect; if A were to go on to write that result back into the database, it would actually leave the database in an inconsistent state. We say that A has seen an inconsistent state of the database and has therefore performed an inconsistent analysis. Note the difference between this example and the previous one: there is no question here of A being dependent on an uncommitted change, since B commits all of its updates before Access Acc 3.

## Locking

A concurrency control technique is called locking. When a transaction needs an assurance that some object it is interested in - typically a database tuple - will not change in some manner while its back is turned, it acquires a lock on that object. The effect of the lock is to "lock other transactions" out of the object in question, and thus, in particular prevent them from changing it.

first, we assume the system supports two kinds of locks, exclusive locks (X-locks) and shared locks (S-locks)

Note: X and S locks are sometimes called write locks and read locks respectively.

- if transaction A holds an exclusive (X-lock) on tuple t, then a request from some distinct transaction B for a lock on either type on t will be denied.
- if transaction A holds on a shared (S-lock) on tuple t, then some
  - A request from<sup>some</sup> distinct transaction B for an X-lock on t will be denied.

- A request from some distinct transaction B for an S-lock on t will be granted (i.e. B will now also hold an S-lock on t).

## ~~Deadlock~~

Deadlock is a situation in which two or more transactions are in a simultaneous wait state, each of them waiting for one of the others to release a lock before it can proceed. Locking can introduce problems of its own, principally the problem of deadlock.

Transaction A	time	Transaction B
—	—	—
LOCK r1 EXCLUSIVE	$t_1$	—
—	—	—
—	$t_2$	LOCK r2 EXCLUSIVE
LOCK r2 EXCLUSIVE	$t_3$	—
wait wait	—	—
wait wait	$t_4$	LOCK r1 EXCLUSIVE
	↓	wait wait

Fig. An example of deadlock

A deadlock. The figure shows a deadlock involving two transactions, but deadlocks involving three, four, ... transactions are also possible, at least in principle; we remark, however, that experiments with system R seemed to show that in practice, deadlocks almost never do involve more than two transactions.

18-18 -> 14

Lab

Customer\_id Customer\_name Address

C001 Smith Ktm

C002 John Bhaktapur

C003 semi Lalitpur

C004 Ivan Ktm

Customer

account\_no balance

A101 900

A102 300

A103 200 account

A104 700

Customer\_id account\_no

C001 A101

C002 A102

C003 A104

C004 A105

depositor

1. Find the name of customer whose customer id is C001
2. Find the name and balance of the customer
3. Insert record to the customer table, customer id: C005, customer\_name: Michael, address: ktm
4. Delete record from depositor whose customer\_id is C004.
5. Increase the balance by 5% in account table whose account no is A101 or current balance is only 200.
6. Find all customer names
7. Find all the different customers address
8. Find the customer\_id and its corresponding customer\_name.
9. List name and address of customer who stay in Kathmandu.
10. List all customer whose name is Smith and address should be ktm
- n. List the customer\_id, acc\_id & balance whose balance is more than 300.  
Select customer.customer\_id, account.account\_no, account.balance from customer, account, depositor where ((customer.